

CS-E4850 Computer Vision, Answers to Exercise Round 9

Yangzhe Kong, Student number: 765756

November 14, 2019

Exercise 1. Neural networks and backpropagation.

1)

$$E(t, y) = \frac{1}{n_2} \sum_i^{n_2} -t_i \log(y_i)$$

2) Let's take a look at each individual z_i

$$E(z_i^{(2)}) = -t_i \log(y_i)$$

Then

$$\begin{aligned} \frac{\partial E}{\partial z_i^{(2)}} &= - \sum_{j=1}^{n_2} t_j \frac{\partial \log(y_j)}{\partial z_i} \\ &= - \sum_{j=1}^{n_2} t_j \frac{1}{y_j} \frac{\partial y_j}{\partial z_i} \\ &= - \frac{t_i}{y_i} \frac{\partial y_i}{\partial z_i} - \sum_{j \neq i}^{n_2} \frac{t_j}{y_j} \frac{\partial y_j}{\partial z_i} \\ &= - \frac{t_i}{y_i} y_i (1 - y_i) - \sum_{j \neq i}^{n_2} \frac{t_j}{y_j} (-y_i y_j) \\ &= -t_i + t_i y_i + \sum_{j \neq i}^{n_2} t_j y_i \\ &= -t_i + \sum_{j=1}^{n_2} t_j y_i \\ &= -t_i + y_i \sum_{j=1}^{n_2} t_j \\ &= y_i - t_i \end{aligned}$$

Thus for the vector $\mathbf{z}^{(2)}$, we have:

$$\frac{\partial E}{\partial \mathbf{z}^{(2)}} = \mathbf{y}^{(1)} - \mathbf{t}^{(1)}$$

3) Still we will prove the formula in the case of individual $y_i^{(1)}$

$$\begin{aligned} \frac{\partial E}{\partial y_i^{(1)}} &= \frac{\partial E(t, s(z_i^{(2)}))}{\partial y_i^{(1)}} \\ &= \frac{\partial E(t, s(y_i^{(1)} \mathbf{W}_i^{(2)}))}{\partial y_i^{(1)}} \\ &= \frac{\partial E(t, s(z_i))}{\partial z_i} \frac{\partial y_i^{(1)} \mathbf{W}_i^{(2)}}{\partial y_i^{(1)}} \\ &= (y_i^{(2)} - t_i)^\top \frac{\partial y_i^{(1)} \mathbf{W}_i^{(2)}}{\partial y_i^{(1)}} \quad (\text{Chain Rule Applied}) \\ &= (y_i^{(2)} - t_i)^\top \mathbf{W}_i^{(2)} \end{aligned}$$

Note that we transpose the first term simply because the derivative of matrix should have the same size of the matrix itself (can also be viewed as to operate the matrix multiplication correctly). This trick also applies to the following formulas.

4)

$$\begin{aligned} \frac{\partial E(t, y_u^{(2)})}{\partial w_{uv}^{(2)}} &= \frac{\partial E(t, s(z_i))}{\partial z_i} \frac{\partial y_v^{(1)} w_{uv}^{(2)}}{\partial w_{uv}^{(2)}} \\ &= (y_u^{(2)} - t_u) \frac{\partial y_v^{(1)} w_{uv}^{(2)}}{\partial w_{uv}^{(2)}} \\ &= (y_u^{(2)} - t_u) y_v^{(1)} \end{aligned}$$

If we write all the results using the formula we derived just now with respect to u and v in a matrix $\nabla \mathbf{W}^2$, we can derive:

$$\begin{aligned} \frac{\partial E}{\partial \mathbf{W}} &= \frac{\partial E(t, s(\mathbf{y}^{(1)} \mathbf{W}))}{\partial \mathbf{W}} \\ &= \frac{\partial E(t, s(\sum_{v=0}^{n^2} \sum_{u=0}^{n^1} y_i^{(1)} w_{uv}^{(2)}))}{\partial w_{uv}} \\ &= (\mathbf{y}^{(2)} - \mathbf{t}) \mathbf{y}^{(1)\top} \end{aligned}$$

5) Since,

$$\frac{\partial \sigma(z)}{\partial z} = \sigma(z)(1 - \sigma(z))$$

We can get:

$$\frac{\partial y^{(1)}}{\partial z^{(1)}} = \frac{\partial \sigma(z^{(1)})}{\partial z^{(1)}} = \sigma(z^{(1)})(1 - \sigma(z^{(1)})) = y^{(1)}(1 - y^{(1)})$$

This can obviously be written as a diagonal matrix as proposed.

6)

$$\begin{aligned} \frac{\partial E}{\partial \mathbf{z}^{(1)}} &= \frac{\partial E(t, s(\mathbf{W}^{(2)}\sigma(\mathbf{z}^{(1)})))}{\partial \mathbf{z}^{(1)}} \\ &= \frac{\partial E(t, s(\mathbf{W}^{(2)}\sigma(\mathbf{z}^{(1)})))}{\partial \mathbf{z}^{(2)}} \frac{\partial \mathbf{z}^{(2)}}{\partial \mathbf{y}^{(1)}} \frac{\partial \mathbf{y}^{(1)}}{\partial \mathbf{z}^{(1)}} \\ &= (\mathbf{y}^{(2)} - \mathbf{t})^\top \mathbf{W}^{(2)} \text{diag}(\mathbf{y}^{(1)}(1 - \mathbf{y}^{(1)})) \end{aligned}$$

7)

$$\begin{aligned} \frac{\partial E}{\partial w_{uv}^{(1)}} &= \frac{\partial E}{\partial z_u^{(1)}} \frac{\partial W_u^{(1)} x_v}{\partial w_{uv}^{(1)}} \\ &= \frac{\partial E}{\partial z_u^{(1)}} x_v \end{aligned}$$

Again we can write the results in the matrix $\nabla \mathbf{W}^{(1)}$

$$\frac{\partial E}{\partial \mathbf{W}^{(1)}} = \left(\frac{\partial E}{\partial \mathbf{z}^{(1)}} \right)^\top \mathbf{x}^\top$$

8) Average the gradients:

$$\begin{aligned} \frac{\partial E}{\partial \mathbf{W}^{(2)}} &= \frac{1}{m} \sum_i^m \frac{\partial E_i}{\partial \mathbf{W}^{(2)}} \\ \frac{\partial E}{\partial \mathbf{W}^{(1)}} &= \frac{1}{m} \sum_i^m \frac{\partial E_i}{\partial \mathbf{W}^{(1)}} \end{aligned}$$

9) $\frac{\partial \frac{\lambda}{2} \|\mathbf{w}\|^2}{\partial \mathbf{w}} = \lambda \mathbf{w}$

Exercise 2. Image classification using a neural network

The code in the function `d_loss_by_d_model` goes as follows: [frame=single]

```

1  ret.input_to_hid = (((model.hid_to_class' * (class_prob - data.
    targets)) .* hid_output .* (1 - hid_output)) * data.inputs
    ') / size(data.targets, 2);
2  ret.hid_to_class = ((class_prob - data.targets) * hid_output') / size(
    data.targets, 2);

```

And here's the resulting training data classification loss:

The total loss on the training data is 2.301907

The classification loss (i.e. without weight decay) on the training data is 2.301907

The classification error rate on the training data is 0.889000

And the plot of training loss and validation loss during training is shown in Figure 1




Figure 1: Training loss and validation loss during training

Exercise 3. Optimisation using backpropagation

We first try with $a2(0, 10, 70, 0.005, 0, false, 4)$, and the resulting training data classification loss is:

The total loss on the training data is 2.301771

The classification loss (i.e. without weight decay) on the training data is 2.301771

The classification error rate on the training data is 0.888000

And the plot of training loss and validation loss during training is shown in Figure 2




Figure 2: Training loss and validation loss during training

We can see that training data loss and validation data loss are both decreasing, but they're still going down steadily after those 70 optimization iterations. We can try with a larger learning rate or with different momentum. And the setting for the experiment goes like this: we try different learning rates 0.002, 0.01, 0.05, 0.2, 1.0, 5.0, and 20.0. We'll try all of those both without momentum (i.e. momentum=0.0 in the program) and with momentum (i.e. momentum=0.9 in the program), so we have a total of $7 \times 2 = 14$ experiments to run. And here are the result:

The learning rate for this experiment is 0.002 and momentum 0

The total loss on the training data is 2.302293

The classification loss (i.e. without weight decay) on the training data is 2.302293

The classification error rate on the training data is 0.897000

The total loss on the validation data is 2.302226

The classification loss (i.e. without weight decay) on the validation data is 2.302226

The classification error rate on the validation data is 0.898000

The total loss on the test data is 2.302253

The classification loss (i.e. without weight decay) on the test data is 2.302253

The classification error rate on the test data is 0.895000

The learning rate for this experiment is 0.002 and momentum 0.9

The total loss on the training data is 2.299567

The classification loss (i.e. without weight decay) on the training data is 2.299567

The classification error rate on the training data is 0.848000

The total loss on the validation data is 2.299502

The classification loss (i.e. without weight decay) on the validation data is 2.299502

The classification error rate on the validation data is 0.839000

The total loss on the test data is 2.299510

The classification loss (i.e. without weight decay) on the test data is 2.299510

The classification error rate on the test data is 0.850778

The learning rate for this experiment is 0.01 and momentum 0

The total loss on the training data is 2.300901

The classification loss (i.e. without weight decay) on the training data is 2.300901

The classification error rate on the training data is 0.886000

The total loss on the validation data is 2.300839

The classification loss (i.e. without weight decay) on the validation data is 2.300839

The classification error rate on the validation data is 0.887000

The total loss on the test data is 2.300854

The classification loss (i.e. without weight decay) on the test data is 2.300854

The classification error rate on the test data is 0.881667

The learning rate for this experiment is 0.01 and momentum 0.9

The total loss on the training data is 2.285062

The classification loss (i.e. without weight decay) on the training data is 2.285062

The classification error rate on the training data is 0.751000

The total loss on the validation data is 2.285036

The classification loss (i.e. without weight decay) on the validation data is 2.285036

The classification error rate on the validation data is 0.756000

The total loss on the test data is 2.284913

The classification loss (i.e. without weight decay) on the test data is 2.284913

The classification error rate on the test data is 0.761667

The learning rate for this experiment is 0.05 and momentum 0

The total loss on the training data is 2.293412

The classification loss (i.e. without weight decay) on the training data is 2.293412

The classification error rate on the training data is 0.765000

The total loss on the validation data is 2.293384

The classification loss (i.e. without weight decay) on the validation data is 2.293384

The classification error rate on the validation data is 0.766000

The total loss on the test data is 2.293329

The classification loss (i.e. without weight decay) on the test data is 2.293329

The classification error rate on the test data is 0.757667

The learning rate for this experiment is 0.05 and momentum 0.9

The total loss on the training data is 1.994052

The classification loss (i.e. without weight decay) on the training data is 1.994052

The classification error rate on the training data is 0.708000

The total loss on the validation data is 1.996559

The classification loss (i.e. without weight decay) on the validation data is 1.996559

The classification error rate on the validation data is 0.700000

The total loss on the test data is 1.989121

The classification loss (i.e. without weight decay) on the test data is 1.989121

The classification error rate on the test data is 0.699889

The learning rate for this experiment is 0.2 and momentum 0

The total loss on the training data is 2.223581

The classification loss (i.e. without weight decay) on the training data is 2.223581

The classification error rate on the training data is 0.801000

The total loss on the validation data is 2.224201

The classification loss (i.e. without weight decay) on the validation data is 2.224201

The classification error rate on the validation data is 0.797000

The total loss on the test data is 2.222849

The classification loss (i.e. without weight decay) on the test data is 2.222849

The classification error rate on the test data is 0.795667

The learning rate for this experiment is 0.2 and momentum 0.9

The total loss on the training data is 1.292496

The classification loss (i.e. without weight decay) on the training data is 1.292496

The classification error rate on the training data is 0.409000

The total loss on the validation data is 1.337265

The classification loss (i.e. without weight decay) on the validation data is 1.337265

The classification error rate on the validation data is 0.423000

The total loss on the test data is 1.301506

The classification loss (i.e. without weight decay) on the test data is 1.301506

The classification error rate on the test data is 0.407000

The learning rate for this experiment is 1 and momentum 0

The total loss on the training data is 1.674933

The classification loss (i.e. without weight decay) on the training data is 1.674933

The classification error rate on the training data is 0.697000

The total loss on the validation data is 1.690025

The classification loss (i.e. without weight decay) on the validation data is 1.690025

The classification error rate on the validation data is 0.714000

The total loss on the test data is 1.676163

The classification loss (i.e. without weight decay) on the test data is 1.676163

The classification error rate on the test data is 0.703222

The learning rate for this experiment is 1 and momentum 0.9

The total loss on the training data is 1.877051

The classification loss (i.e. without weight decay) on the training data is 1.877051

The classification error rate on the training data is 0.721000

The total loss on the validation data is 1.920913

The classification loss (i.e. without weight decay) on the validation data is 1.920913

The classification error rate on the validation data is 0.737000

The total loss on the test data is 1.897451

The classification loss (i.e. without weight decay) on the test data is 1.897451

The classification error rate on the test data is 0.724222

The learning rate for this experiment is 5 and momentum 0

The total loss on the training data is 2.301619

The classification loss (i.e. without weight decay) on the training data is 2.301619

The classification error rate on the training data is 0.899000

The total loss on the validation data is 2.302043

The classification loss (i.e. without weight decay) on the validation data is 2.302043

The classification error rate on the validation data is 0.900000

total loss on the test data is 2.302443

The classification loss (i.e. without weight decay) on the test data is 2.302443

The classification error rate on the test data is 0.899333

The learning rate for this experiment is 5 and momentum 0.9

The total loss on the training data is 2.302585

The classification loss (i.e. without weight decay) on the training data is 2.302585

The classification error rate on the training data is 0.886000

The total loss on the validation data is 2.302585

The classification loss (i.e. without weight decay) on the validation data is 2.302585

The classification error rate on the validation data is 0.897000

The total loss on the test data is 2.302585

The classification loss (i.e. without weight decay) on the test data is 2.302585

The classification error rate on the test data is 0.887778

The learning rate for this experiment is 20 and momentum 0

The total loss on the training data is 2.302585

The classification loss (i.e. without weight decay) on the training data is 2.302585

The classification error rate on the training data is 0.900000

The total loss on the validation data is 2.302585

The classification loss (i.e. without weight decay) on the validation data is 2.302585

The classification error rate on the validation data is 0.900000

The total loss on the test data is 2.302585

The classification loss (i.e. without weight decay) on the test data is 2.302585

The classification error rate on the test data is 0.900000

The learning rate for this experiment is 20 and momentum 0.9

The total loss on the training data is 2.302585

The classification loss (i.e. without weight decay) on the training data is 2.302585

The classification error rate on the training data is 0.874000

The total loss on the validation data is 2.302585

The classification loss (i.e. without weight decay) on the validation data is 2.302585

The classification error rate on the validation data is 0.882000

The total loss on the test data is 2.302585

The classification loss (i.e. without weight decay) on the test data is 2.302585

The classification error rate on the test data is 0.877889

And all training set and validation set results are shown in the following table:

Learning rate	Without momentum				With momentum 0.9			
	training loss	training error	validation loss	validation error	training loss	training error	validation loss	validation error
0.002	2.302	0.897	2.302	0.898	2.300	0.848	2.300	0.839
0.01	2.301	0.886	2.301	0.887	2.285	0.751	2.285	0.756
0.05	2.293	0.765	2.293	0.766	1.994	0.708	1.997	0.700
0.2	2.224	0.801	2.224	0.797	1.292	0.409	1.337	0.423
1.0	1.675	0.697	1.690	0.714	1.877	0.721	1.921	0.737
5.0	2.302	0.899	2.302	0.900	2.303	0.886	2.303	0.897
20.0	2.303	0.900	2.303	0.900	2.303	0.874	2.303	0.882

We can observe from the results that if we choose a reasonable learning rate (0.2) and implement with momentum, we can achieve a overall better result. If we use extreme small or big learning rate, the performance will be worse. But one thing worth noting is that although small learning rate may perform worse compared with larger learning rate under the same number of iterations, if we train for larger iterations, small learning rate may still able to converge. Thus, number of training iterations is also an important parameter. So, in practice even if we build a very good model, we still have to tune our hyperparameters to train it efficiently to achieve best performance.

Finally we can train our model for larger iterations and see what we can get with $a2(0, 10, 200, 0.2, 0.9, false, 4)$. The plot is shown in Figure 3




Figure 3: The training loss and validation loss during training with $a2(0, 10, 200, 0.2, 0.9, false, 4)$

We can see that at iteration 140 the model achieves best performance.